

1R Mushroom Example

Name: Perry Koob

eMail: koobp@mst.edu

Course: CS 5402

Date: 2020-06-21

Course: Version 1.3.2

Concept Description:


Train a system from existing data to classify mushrooms as either edible or inedible.

Data Collection:

The data has been provided by Perry B. Koob, not professor or doctor. It is a modified version of the UCI Mushroom data set found here:

<https://archive.ics.uci.edu/ml/datasets/mushroom>

Example Description:



Mushroom Diagram

edible.poisonous

This is the class label.

cap.shape

Nominal attribute that describes the cap shape of the mushroom as:

bell
conical convex
flat
knobbed sunken

cap.surface

Nominal attribute that describes the cap surface of the mushroom as:

fibrous grooves scaly
smooth

cap.color

Nominal attribute that describes the cap color of the mushroom as:

brown
buff
cinnamon green
grey
pink
purple
red
white
yellow

bruises

Nominal attribute, boolean in nature, that describes if the mushroom has bruises.

odor

Nominal attribute that describes the odor of the mushroom as:

almond
anise
creosote fishy
foul
musty
none
pungent spicy

colony

Interval attribute that describes the approximate size of the mushroom colony:

edible

This is a binary class label generated from the edible.poisonous class label. It is a transformation of a nominal Class label, so it is also nominal. The labels are now edible or inedible.

There are no missing values.

Data Import and Wrangling:

```
df <- read.csv(file = '../src-data/mushroom.csv', stringsAsFactors=TRUE)
```

Create a new Class in the training dataset

```
df$edible <- ifelse(df$edible.poisonous == 'edible', "edible", "inedible")
df$edible <- as.factor(df$edible)
```

Partition the data into a training set and a test set using a 75/25 split.

```
set.seed(123)
size = floor(0.75*nrow(df))

train_index = sample(seq_len(nrow(df)), size = size)

train = df[train_index,]
test=df[-train_index,]
```

Exploratory Data Analysis:

Looking into what type of measure the attributes are.

```
summary(df)
```


##	edible.poisonous	cap.shape	cap.surface	cap.color	bruises
##	edible :4208	bell : 452	fibrous:2320	brown :2284	
	bruises:3376				
##	poisonous:3916	conical: 4	grooves: 4	grey :1840	no
	:4748				
##		convex :3656	scaly :3244	red :1500	
##		flat :3152	smooth :2556	yellow :1072	
##		knobbed: 828		white :1040	
##		sunken : 32		buff : 168	
##				(Other): 220	
##	odor	colony	edible		
##	none :3528	Min. : 1.00	edible :4208		
##	foul :2160	1st Qu.: 1.00	inedible:3916		
##	fishy : 576	Median : 2.00			
##	spicy : 576	Mean : 46.47			
##	almond : 400	3rd Qu.: 7.00			

```
## anise : 400 Max. :999.00  
## (Other): 484
```

```
#Cap Shape
```

```
count <- table(train$cap.shape)
```

```
barplot(count,main="Count of Cap Shape", xlab="Cap Shape")
```

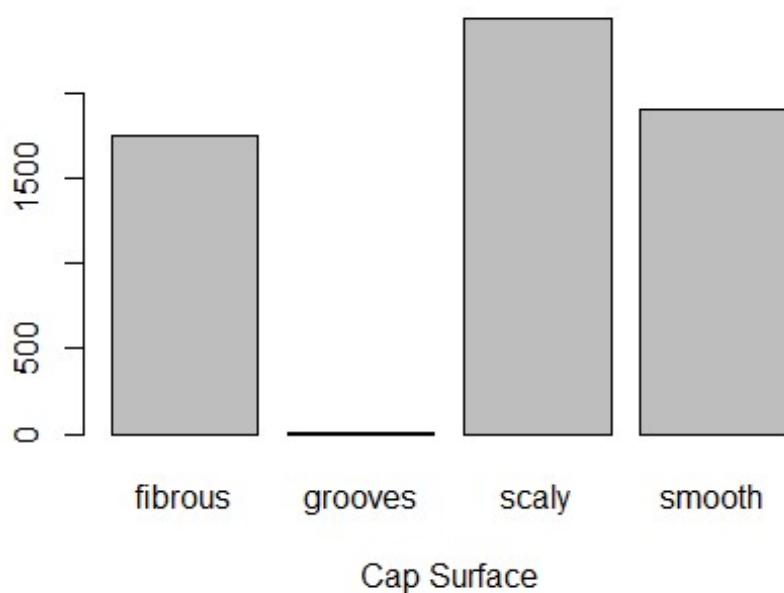


```
#Cap Surface
```

```
count <- table(train$cap.surface)
```

```
barplot(count,main="Count of Cap Surface Sizes", xlab="Cap Surface")
```

Count of Cap Surface Sizes

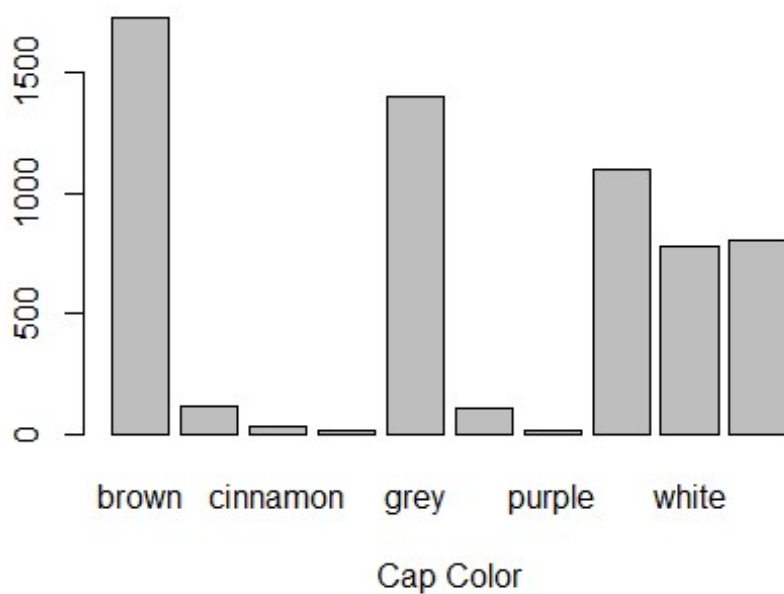


```
#Cap Color
```


```
count <- table (train$cap.color)
```

```
barplot(count,main="Count of Cap Color", xlab="Cap Color")
```

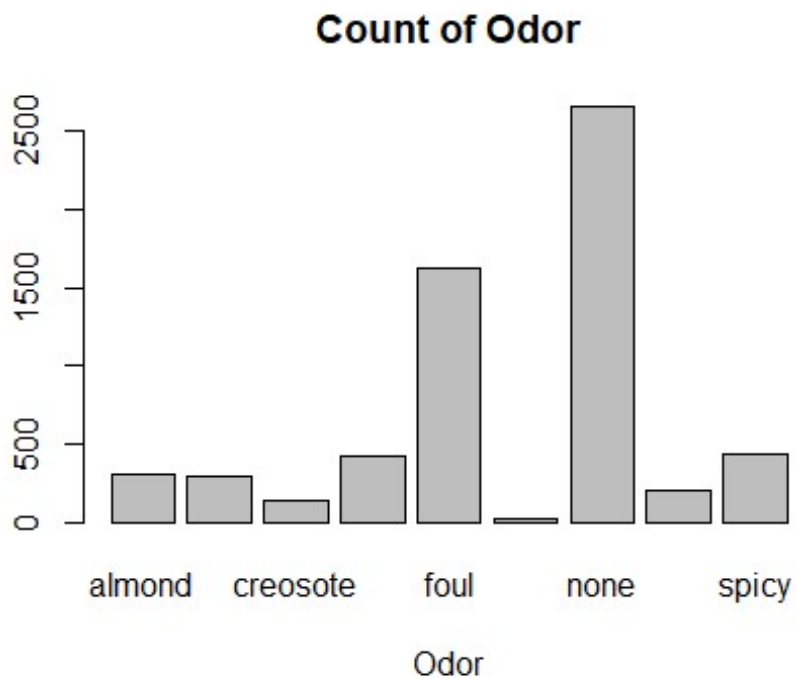
Count of Cap Color



```
#Bruises  
count <- table (train$bruises)  
barplot(count,main="Count of Bruises", xlab="Bruises")
```



```
#Odor  
count <- table (train$odor)  
barplot(count,main="Count of Odor", xlab="Odor")
```



#Determine what the missing attribute label is for odor

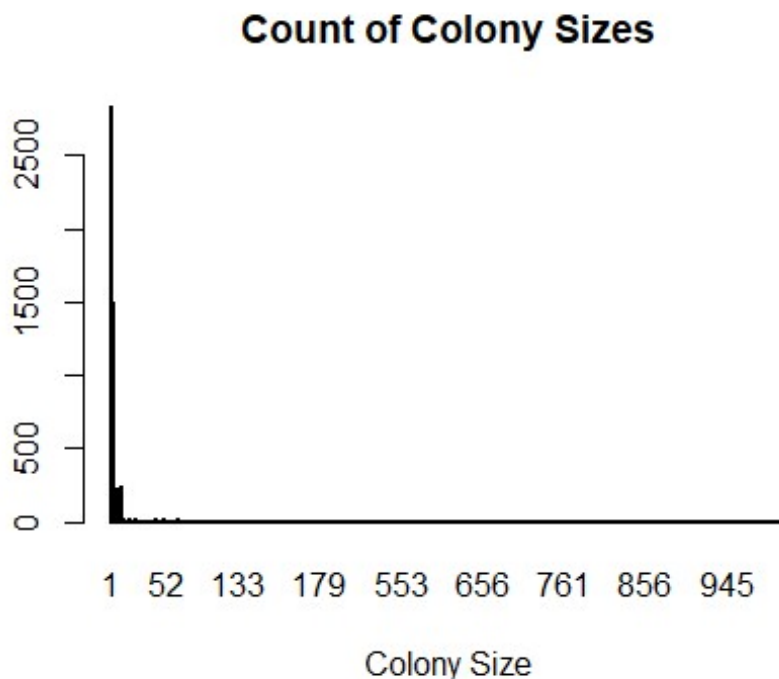
```
summary(train$odor)
```

```
##  almond  anise creosote  fishy  foul  musty  none  pungent
##    299    293    140    427  1628    20  2659    197
##  spicy
##    430
```

#Colony

```
count <- table (train$colony)
```

```
barplot(count,main="Count of Colony Sizes", xlab="Colony Size")
```



Looking at the bar chart we can see Colony Size is not nominal. Since we need nominal values to be able to run 1R, we are going to need to discretize the Colony Size.

```
str(train$colony)
##  int [1:6093] 1 1 855 8 1 2 1 6 1 1 ...
```

Discretizing the Colony attribute

Discretizing Using the Holte method

```
# Holte discretization has not been updated for R version 4
#disc.1r(df)
```

Discretizing Using frequency

Since the Holte method of Discretization is not working for this version, we will use another discretization method. The Frequency based method.

```
discretize(train$colony, onlycuts = TRUE)

## Warning in discretize(train$colony, onlycuts = TRUE): The calculated
breaks are: 1, 1, 2, 999
## Only unique breaks are used reducing the number of intervals. Look at ?
discretize for details.

## [1] 1 2 999

train$dcolony[train$colony <=2] <- "single"
train$dcolony[train$colony > 2] <- "colony"
```



```

train$dcolony <- as.factor(train$dcolony)

test$dcolony[test$colony <=2] <- "single"
test$dcolony[test$colony > 2] <- "colony"
test$dcolony <- as.factor(test$dcolony)

```

Mining or Analytics:

Getting the 1R rules using a manual method.

```

# Initialize prediction data frame with the train data frame.
prediction <- train

```

Cap Shape

```

# Get the frequency of the Cap.shapes and their frequency by classes values
cap.shape.count<- as.data.frame(table(train$cap.shape))
cap.shape.true <- as.data.frame(table(train[train$edible ==
"edible",]$cap.shape))
cap.shape.false <- as.data.frame(table(train[train$edible ==
"inedible",]$cap.shape))

# Rename the columns so the make sense after the merge
colnames(cap.shape.count) <- c("Label","Freq")
colnames(cap.shape.true) <- c("Label","Edible")
colnames(cap.shape.false) <- c("Label","Inedible")

# Merge the three dataframes together
cap.shape <- merge(cap.shape.count, cap.shape.true, by="Label", all=TRUE)
cap.shape <- merge(cap.shape, cap.shape.false, by="Label", all=TRUE)
cap.shape

##      Label Freq Edible Inedible
## 1    bell   342     310        32
## 2 conical    3        0         3
## 3  convex 2721    1447    1274
## 4   flat 2388    1205    1183
## 5 knobbed   614     170     444
## 6  sunken    25      25         0

```

Base on the frequency of edible and inedible, we would generate the following rule set for the Cap Shape attribute:

```

bell -> Edible
conical -> Inedible
convex -> Edible
flat -> Edible
knobbed -> Inedible
sunken -> Edible

```

Now we predict what the Cap Shape rules would classify our training data as.

```

prediction$pred.cap.shape <- NA

prediction$pred.cap.shape[prediction$cap.shape == "bell"] <- "edible"
prediction$pred.cap.shape[prediction$cap.shape == "conical"] <- "inedible"
prediction$pred.cap.shape[prediction$cap.shape == "convex"] <- "edible"
prediction$pred.cap.shape[prediction$cap.shape == "flat"] <- "edible"
prediction$pred.cap.shape[prediction$cap.shape == "knobbed"] <- "inedible"
prediction$pred.cap.shape[prediction$cap.shape == "sunken"] <- "edible"

prediction$cap.shape <- NULL

prediction$pred.cap.shape <- as.factor(prediction$pred.cap.shape)

```

Cap Surface

```

cap.surface.count<- as.data.frame(table(train$cap.surface))
cap.surface.true <- as.data.frame(table(train[train$edible ==
"edible",]$cap.surface))
cap.surface.false <- as.data.frame(table(train[train$edible ==
"inedible",]$cap.surface))

colnames(cap.surface.count) <- c("Label", "Freq")
colnames(cap.surface.true) <- c("Label", "Edible")
colnames(cap.surface.false) <- c("Label", "Inedible")

cap.surface <- merge(cap.surface.count, cap.surface.true, by="Label", all=TRUE)
cap.surface <- merge(cap.surface, cap.surface.false, by="Label", all=TRUE)
cap.surface

##      Label Freq Edible Inedible
## 1 fibrous 1754   1196     558
## 2 grooves    3     0        3
## 3 scaly    2434   1112    1322
## 4 smooth   1902    849    1053

```

Base on the frequency of edible and inedible, we would generate the following rule set for the Cap Surface attribute:

```

fibrous -> Edible
grooves -> Inedible
scaly -> Inedible
smooth -> Inedible

```

Now we predict what the Cap Surface rules would classify our training data as.

```

prediction$pred.cap.surface <- NA

prediction$pred.cap.surface[prediction$cap.surface == "fibrous"] <- "edible"
prediction$pred.cap.surface[prediction$cap.surface == "grooves"] <-

```

```

"inedible"
prediction$pred.cap.surface[prediction$cap.surface == "scaly"] <- "inedible"
prediction$pred.cap.surface[prediction$cap.surface == "smooth"] <- "inedible"

prediction$cap.surface <- NULL

prediction$pred.cap.surface <- as.factor(prediction$pred.cap.surface)

```

Cap Color

```

cap.color.count<- as.data.frame(table(train$cap.color))
cap.color.true <- as.data.frame(table(train[train$edible ==
"edible",]$cap.color))
cap.color.false <- as.data.frame(table(train[train$edible ==
"inedible",]$cap.color))

colnames(cap.color.count) <- c("Label","Freq")
colnames(cap.color.true) <- c("Label","Edible")
colnames(cap.color.false) <- c("Label","Inedible")

cap.color <- merge(cap.color.count, cap.color.true, by="Label", all=TRUE)
cap.color <- merge(cap.color, cap.color.false, by="Label", all=TRUE)
cap.color

```

	Label	Freq	Edible	Inedible
## 1	brown	1729	945	784
## 2	buff	118	34	84
## 3	cinnamon	32	25	7
## 4	green	16	16	0
## 5	grey	1400	788	612
## 6	pink	106	40	66
## 7	purple	9	9	0
## 8	red	1098	454	644
## 9	white	781	550	231
## 10	yellow	804	296	508

Base on the frequency of edible and inedible, we would generate the following rule set for the Cap Color attribute:

```

brown -> Edible
buff -> Inedible
cinnamon -> Edible
green -> Edible
grey -> Edible
pink -> Inedible
purple -> Edible
red -> Inedible

```

white -> Edible
yellow -> Inedible

Now we predict what the Cap Color rules would classify our training data as.

```
prediction$pred.cap.color <- NA

prediction$pred.cap.color[prediction$cap.color=="brown"] <- "edible"
prediction$pred.cap.color[prediction$cap.color=="buff"] <- "inedible"
prediction$pred.cap.color[prediction$cap.color=="cinnamon"] <- "edible"
prediction$pred.cap.color[prediction$cap.color=="green"] <- "edible"
prediction$pred.cap.color[prediction$cap.color=="grey"] <- "edible"
prediction$pred.cap.color[prediction$cap.color=="pink"] <- "inedible"
prediction$pred.cap.color[prediction$cap.color=="purple"] <- "edible"
prediction$pred.cap.color[prediction$cap.color=="red"] <- "inedible"
prediction$pred.cap.color[prediction$cap.color=="white"] <- "edible"
prediction$pred.cap.color[prediction$cap.color=="yellow"] <- "inedible"

prediction$cap.color <- NULL
prediction$pred.cap.color <- as.factor(prediction$pred.cap.color)
```

Bruises

```
bruises.count<- as.data.frame(table(train$bruises))
bruises.true <- as.data.frame(table(train[train$edible ==
"edible",]$bruises))
bruises.false <- as.data.frame(table(train[train$edible ==
"inedible",]$bruises))

colnames(bruises.count) <- c("Label","Freq")
colnames(bruises.true) <- c("Label","Edible")
colnames(bruises.false) <- c("Label","Inedible")

bruises <- merge(bruises.count,bruises.true, by="Label",all=TRUE)
bruises <- merge(bruises,bruises.false, by="Label",all=TRUE)
bruises

##      Label Freq Edible Inedible
## 1 bruises 2506   2044      462
## 2      no 3587   1113     2474
```

Base on the frequency of edible and inedible, we would generate the following rule set for the Cap Color attribute:

bruises -> Edible
no -> Inedible

Now we predict what the Bruises rules would classify our training data as.

```

prediction$pred.bruises <- NA

prediction$pred.bruises[prediction$bruises=="bruises"] <- "edible"
prediction$pred.bruises[prediction$bruises=="no"] <- "inedible"

prediction$bruises <- NULL
prediction$pred.bruises <- as.factor(prediction$pred.bruises)

```

Odor

```

odor.count<- as.data.frame(table(train$odor))
odor.true <- as.data.frame(table(train[train$edible == "edible",]$odor))
odor.false <- as.data.frame(table(train[train$edible == "inedible",]$odor))

colnames(odor.count) <- c("Label","Freq")
colnames(odor.true) <- c("Label","Edible")
colnames(odor.false) <- c("Label","Inedible")

odor <- merge(odor.count,odor.true, by="Label",all=TRUE)
odor <- merge(odor,odor.false, by="Label",all=TRUE)
odor

```

##	Label	Freq	Edible	Inedible
## 1	almond	299	299	0
## 2	anise	293	293	0
## 3	creosote	140	0	140
## 4	fishy	427	0	427
## 5	foul	1628	0	1628
## 6	musty	20	0	20
## 7	none	2659	2565	94
## 8	pungent	197	0	197
## 9	spicy	430	0	430

Base on the frequency of edible and inedible, we would generate the following rule set for the Odor attribute:

```

almond -> edible
anise -> edible
creosote -> inedible
fishy -> inedible
foul -> inedible
musty -> inedible
none -> edible
pungent -> inedible
spicy -> inedible

```

Now we predict what the Odor rules would classify our training data as.

```

prediction$pred.odor <- NA

prediction$pred.odor[prediction$odor=="almond"] <- "edible"
prediction$pred.odor[prediction$odor=="anise"] <- "edible"
prediction$pred.odor[prediction$odor=="creosote"] <- "inedible"
prediction$pred.odor[prediction$odor=="fishy"] <- "inedible"
prediction$pred.odor[prediction$odor=="foul"] <- "inedible"
prediction$pred.odor[prediction$odor=="musty"] <- "inedible"
prediction$pred.odor[prediction$odor=="none"] <- "edible"
prediction$pred.odor[prediction$odor=="pungent"] <- "inedible"
prediction$pred.odor[prediction$odor=="spicy"] <- "inedible"

prediction$odor <- NULL
prediction$pred.odor <- as.factor(prediction$pred.odor)

```

The error rate will be determined for each Rule Set to determine which one is better.

```

# Determine when the rule predict incorrectly
prediction$pred.cap.shape.correct = ifelse(prediction$edible ==
prediction$pred.cap.shape,0,1)
prediction$pred.cap.surface.correct = ifelse(prediction$edible ==
prediction$pred.cap.surface,0,1)
prediction$pred.cap.color.correct = ifelse(prediction$edible ==
prediction$pred.cap.color,0,1)
prediction$pred.bruises.correct = ifelse(prediction$edible ==
prediction$pred.bruises,0,1)
prediction$pred.odor.correct = ifelse(prediction$edible ==
prediction$pred.odor,0,1)

# Calculate the error rate
print(paste("Error Rate of Cap Shape
Rules",sum(prediction$pred.cap.shape.correct)/nrow(prediction)))

## [1] "Error Rate of Cap Shape Rules 0.436402429016905"

print(paste("Error Rate of Cap Surface
Rules",sum(prediction$pred.cap.surface.correct) /nrow(prediction)))

## [1] "Error Rate of Cap Surface Rules 0.413425242081077"

print(paste("Error Rate of Cap Color
Rules",sum(prediction$pred.cap.color.correct)/nrow(prediction)))

## [1] "Error Rate of Cap Color Rules 0.403413753487609"

print(paste("Error Rate of Bruises
Rules",sum(prediction$pred.bruises.correct)/nrow(prediction)))

## [1] "Error Rate of Bruises Rules 0.258493353028065"

print(paste("Error Rate of Cap Odor
Rules",sum(prediction$pred.odor.correct)/nrow(prediction)))

```

```
## [1] "Error Rate of Cap Odor Rules 0.0154275397997702"
```

The Rule Set with the best error rate is the Odor Rules Set, which has an error rate of 0.015.

So we choose the the Odor Rules Set as our One Rule.

Getting the 1R rules using the OneR package.

```
df_train <- train
df_train$edible.poisonous <- NULL

model <- OneR(df_train)

## Warning in OneR.data.frame(df_train): data contains unused factor levels

model_prediction <- predict(model, test)
```

Evaluation:

We can evaluate the health of the two models using the training data set and Confusion Matrices obtained from both models. The Accuracy rate and the Error rate will allow us to compare the models to each other.

```
#prepare the list of classes from the test data for evaluation
reference <- as.data.frame(test$edible)
colnames(reference) <- c("class")
reference <- as.factor(reference$class)

#prepare the list of predictions from the test data for evaluation
testing <- as.data.frame(test[c("odor")])
colnames(testing) <- c("odor")

testing$pred[testing$odor=="almond"] <- "edible"
testing$pred[testing$odor=="anise"] <- "edible"
testing$pred[testing$odor=="creosote"] <- "inedible"
testing$pred[testing$odor=="fishy"] <- "inedible"
testing$pred[testing$odor=="foul"] <- "inedible"
testing$pred[testing$odor=="musty"] <- "inedible"
testing$pred[testing$odor=="none"] <- "edible"
testing$pred[testing$odor=="pungent"] <- "inedible"
testing$pred[testing$odor=="spicy"] <- "inedible"

testing$pred <- as.factor(testing$pred)

confusionMatrix(testing$pred, reference)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction edible inedible
## edible      1051      26
```

```

##      inedible      0      954
##
##              Accuracy : 0.9872
##              95% CI : (0.9813, 0.9916)
##      No Information Rate : 0.5175
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9743
##
##      McNemar's Test P-Value : 9.443e-07
##
##              Sensitivity : 1.0000
##              Specificity : 0.9735
##              Pos Pred Value : 0.9759
##              Neg Pred Value : 1.0000
##              Prevalence : 0.5175
##              Detection Rate : 0.5175
##      Detection Prevalence : 0.5303
##              Balanced Accuracy : 0.9867
##
##              'Positive' Class : edible
##

```

```
eval_model(model_prediction, test)
```

```

##
## Confusion matrix (absolute):
##      Actual
## Prediction colony single Sum
##      colony      259      88 347
##      single      352     1332 1684
##      Sum          611     1420 2031
##
## Confusion matrix (relative):
##      Actual
## Prediction colony single Sum
##      colony      0.13      0.04 0.17
##      single      0.17      0.66 0.83
##      Sum          0.30      0.70 1.00
##
## Accuracy:
## 0.7834 (1591/2031)
##
## Error rate:
## 0.2166 (440/2031)
##
## Error rate reduction (vs. base rate):
## 0.2799 (p-value < 2.2e-16)

```


The accuracy rate of the 1R Rules generated by hand is 98% and generated by the oneR package is 78%. So both are good models.

Results:

Since both the hand generated 1R classifier and the oneR package generated the same results, a 1R classifier based on the Odor attribute, and they both have a acceptable accuracy with training and test data, we can feel confident that our model is a strong one.

Our final model is...

If the odor of the mushroom is:

almond -> edible
anise -> edible
creosote -> inedible
fishy -> inedible
foul -> inedible
musty -> inedible
none -> edible
pungent -> inedible
spicy -> inedible

References:

<https://stackoverflow.com/questions/7739578/merge-data-frames-based-on-rownames-in-r> <https://www.datanovia.com/en/lessons/rename-data-frame-columns-in-r/>
<https://www.dummies.com/programming/r/how-to-convert-tables-to-a-data-frame-in-r/>
https://wilkelab.org/classes/SDS348/2016_spring/labs/remove_variables.html
<https://stackoverflow.com/questions/2854625/select-only-rows-if-its-value-in-a-particular-column-is-less-than-its-value-in-t>
<https://stackoverflow.com/questions/7072159/how-do-you-remove-columns-from-a-data-frame> <https://stats.stackexchange.com/questions/5253/how-do-i-get-the-number-of-rows-of-a-data-frame-in-r>
<https://www.rdocumentation.org/packages/caret/versions/3.45>
<https://www.rdocumentation.org/packages/arules/versions/1.6-8/topics/discretize>